

VRDL HW2: Digit Recognition Report

Yi-Hsiang Ho, 111550106

1. Introduction

This task is to locate the digit in an image and recognize it using Faster R-CNN. The core idea of this work is to leverage restart learning rate scheduler and adjust anchor aspect and size. ResNet-101 is chosen to be the backbone of the model. GitHub repository is available [here](#).

1.1. Faster R-CNN

Faster R-CNN is a widely used deep learning framework for object detection. It builds upon earlier R-CNN models by integrating region proposal generation and object classification into a single network, significantly improving both speed and accuracy. The architecture consists of a convolutional backbone for feature extraction, a Region Proposal Network (RPN) that suggests potential object locations, and a head inherited from Fast R-CNN that refines these proposals and classifies the objects. Unlike its predecessors, Faster R-CNN eliminates the need for external region proposal methods, such as selective search, by using the RPN to efficiently generate high-quality proposals directly from the feature maps. [2]

1.2. Region Proposal Network (RPN)

The Region Proposal Network (RPN) is a key component of Faster R-CNN that efficiently generates object proposals directly from feature maps produced by the convolutional backbone. RPN slides a small network over the feature map and predicts objectness scores and bounding box for multiple anchors. These anchors vary in scale and aspect ratio, allowing the RPN to detect objects of different shapes and sizes. By sharing convolutional features with the detection network, the RPN significantly speeds up the proposal process while maintaining high accuracy.

1.3. Restart Learning Rate Scheduler

Restart is a technique used in training neural networks to avoid local minima by periodically resetting the learning rate. Instead of steadily decreasing the learning rate throughout training, this approach reduces it over time but then abruptly increases it at predefined intervals, effectively restarting the optimization process. This strategy encourages the model to explore new regions of the loss landscape after each restart, often leading to better generalization.

2. Method

2.1. Model Architecture

The backbone was changed from ResNet-50 (default) to ResNet-101 for a better feature extraction. It is trained from the ImageNet pretrained Weights retrieved from PyTorch as `ResNet101_Weights.IMAGENET1K_V2`.

The architecture of RPN doesn't change. But the anchor size and aspect ratio are adjusted to fit this task. Since digits are usually small and rectangular, the anchor size is set from 16x16 to 256x256 and the aspect ratio is set to 1:4, 1:2, 1:1, 2:1, and 4:1. There are 25 anchors in total. This allows the model to better capture the bounding boxes of digits. [1]

After region proposals are generated and refined (by NMS) by RPN, they are projected onto the shared feature map and passed through a RoI pooling to produce fixed-size feature vectors. These vectors are then fed into 2 layers of MLP and 1 branched MLP to output classification and bounding box regression. This part keeps the same as the initial implementation.

2.2. Optimization

For optimization, AdamW is used with an initial learning rate of 5e-5. A cosine annealing scheduler dynamically adjusts the learning rate throughout training. A restart strategy is adopted to avoid local minima. It will restart after the first 10,000 steps, and the number of steps between each restart will double. That is, the next restart will occur after the next 20,000 steps, which is at step 30,000. Since the total number of steps is about 70,000, it will restart three times. Fig. 1 visualizes the change of the learning rate.

Training process is performed with a batch size of 8 over 20 epochs. The provided dataset contains roughly 30,000 training images, 3,000 validation images, and 13,000 test images. The best-performing model is selected based on the validation accuracy. The model is trained on a single NVIDIA GeForce RTX 4090 GPU in about 8.5 hours.

2.3. Inference

During inference, an image will be fed into the model to generate labels, bounding boxes and the confidence score. The boxes with score lower than 0.5 will be discarded. The remaining boxes will be involved in the final number prediction. The final prediction is based on the location of the bounding box. Sort the boxes in one image from left to

right, top to bottom, and concatenate the label for the sorted boxes. The final result is a string of digits like "123".

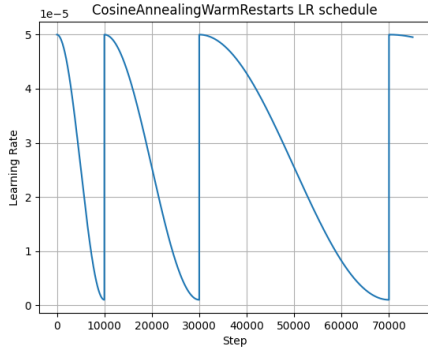


Figure 1. **Visualization of restart learning rate scheduler.** The learning rate is reduced over time but then abruptly increased at predefined intervals.

2.4. Hyperparameters

The hyperparameter details are listed below:

- Learning rate: 5e-5
- Optimizer: AdamW
- Scheduler: CosineAnnealingWarmRestarts ($T_0 = 10,000$, $T_{mul} = 2$)
- Batch Size: 16
- Epochs: 20
- Anchor size: 16x16, 32x32, 64x64, 128x128, 256x256
- Anchor aspect: 0.25 (1:4), 0.5 (1:2), 1.0 (1:1), 2.0 (2:1), 4.0 (4:1)

3. Results

In this section, I compare the performance of each component in the model. The details of each method are listed:

- Res101: ResNet 101 pretrained on ImageNet without any tricks.
- Restart: Res101 with restart learning rate scheduler.
- Anchor: Change anchor size and aspect ratio.
- All: Combines all the techniques mentioned in Sec. 2.

The validation accuracy is calculated by comparing the final result of the model and the ground truth for each image. The ground truth labels and boxes are converted to a string of digits using the same method mentioned above. While the mAP is calculated per boxes and labels between the predicted and ground truth. The mAP is calculated as average AP at each IoU threshold, from 0.5 to 0.95 with a step of 0.05, and each label.

The results are shown in Tab. 1 and Tab. 2. All the results here are evaluated by choosing the best checkpoint based on the validation accuracy. Using customized anchors shows a great improvement in the accuracy while the restart strategy improves the mAP much. The combination of all the techniques achieves the best performance.

Method	Val	Test pub.	Test priv.
Res101	0.8063	0.7536	0.7557
Restart	0.8045	0.7580	0.7564
Anchor	0.8168	0.7777	0.7727
All	0.8177	0.7840	0.7810

Table 1. **Accuracy results of different models.** "Val" refers to validation accuracy. "Test pub." and "Test priv." refer to public and private test set accuracy, respectively. The highest values in each column are highlighted in bold.

Method	Val	Test pub.	Test priv.
Res101	0.4526	0.3686	0.3681
Restart	0.4554	0.3770	0.3777
Anchor	0.4524	0.3720	0.3706
All	0.4581	0.3751	0.3777

Table 2. **mAP results of different models.** The definition of the term is the same as Tab. 1. The highest values in each column are highlighted in bold.

The validation accuracy, mAP and training loss curve are shown in Fig. 2, Fig. 3 and Fig. 4, respectively. At first, I found that Res101 has a very steady training process. The accuracy and mAP keep increasing, the loss keeps decreasing, and it converges quickly. Thus, I plan to use restart strategy to add some interference in the training process. The restart strategy works well. While a restart happens (roughly at epoch 2, 7, and 18), the accuracy and mAP drop a little but then increase rapidly. The training loss also rises after each restart. This indicates that the model is exploring new regions of the loss landscape.




Figure 2. **Validation accuracy curve.**




Figure 3. Validation mAP curve.




Figure 4. Training loss curve.

Other Experiments

The Selection of Best Checkpoint

The best checkpoint is selected based on the validation accuracy. I’ve also tried to select based on the mAP. However, the mAP is not a good indicator of the model’s performance. The mAP is calculated based on the predicted boxes and labels. However, the predicted boxes are not always accurate. The mAP is not sensitive to the accuracy of the predicted boxes. Therefore, I choose to select the best checkpoint based on the validation accuracy. Tab. 3 verify this thought.

Method	Accuracy			mAP		
	Val	Pub.	Priv.	Val	Pub.	Priv.
Res101 acc.	0.806	0.754	0.756	0.453	0.369	0.368
Res101 mAP	0.792	0.722	0.720	0.459	0.367	0.370
Restart acc.	0.807	0.767	0.765	0.455	0.377	0.378
Restart mAP	0.805	0.758	0.756	0.461	0.376	0.377

Table 3. **The result using different selection method.** "acc." means the result of selecting the best checkpoint based on the validation accuracy, while "mAP" selects based on the mAP. The higher values are highlighted in bold. "Pub" is the same as "Test pub." in previous section, as well as "Priv."

The Selection of Backbone

The selection of the convolutional backbone may affect the model’s performance. I’ve conducted experiments with different backbones, ResNet50, ResNet101, ResNeXt50 and ResNeXt101 [3]. The results are shown in Tab. 4. Both ResNet101 and ResNeXt101 perform well, but ResNeXt101 has a poor performance when integrating restart and anchor trick. It also needs 1.7x longer time to train than ResNet101. Therefore, I choose ResNet101 as the backbone of the model.

Method	Accuracy			mAP		
	Val	Pub.	Priv.	Val	Pub.	Priv.
ResNet50	0.675	0.661	0.657	0.426	0.361	0.360
ResNet101	0.806	0.754	0.756	0.453	0.369	0.368
ResNeXt50	0.746	0.680	0.690	0.448	0.356	0.359
ResNeXt101	0.808	0.757	0.761	0.450	0.365	0.366

Table 4. **The result of different backbone.** The highest values for each column are highlighted in bold.

Sharpening

I’ve also experimented with sharpening the input image. I found that the input image may be too blurred to recognize even by human. Therefore, I tried to sharpen the image before passing it to the model. It is implemented by PIL package. However it does not consistently improve performance. The results are shown in Tab. 5. Thus, it is not included in the final model.

Method	Accuracy			mAP		
	Val	Pub.	Priv.	Val	Pub.	Priv.
ResNet101	0.806	0.754	0.756	0.453	0.369	0.368
ResNet101*	0.796	0.733	0.740	0.457	0.367	0.367
Restart	0.807	0.767	0.765	0.455	0.377	0.378
Restart*	0.808	0.770	0.771	0.457	0.373	0.377
All	0.818	0.784	0.781	0.458	0.375	0.378
All*	0.814	0.786	0.786	0.451	0.374	0.371

Table 5. **The result of adding sharpening.** The higher values for each column are highlighted in bold. * means adding sharpening.

References

- [1] Object detection finetuning tutorial - pytorch tutorials 2.6.0+cu124 documentation. 1
- [2] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. Faster r-cnn: Towards real-time object detection with region proposal networks. In *NeurIPS*, 2015. 1

- [3] Saining Xie, Ross Girshick, Piotr Dollár, Zhuowen Tu, and Kaiming He. Aggregated residual transformations for deep neural networks. In *CVPR*, 2017. 3